

Conditional Drifting Models for Learned Channel Simulation:

Detailed Experimental Draft and Comparison With Diffusion and GAN Baselines

Anonymous Submission
DM_for_learning_channels

Abstract—Accurate stochastic channel simulation is central to data-driven communication system design, especially when measured channels are costly to acquire repeatedly or when differentiable learned surrogates are needed inside end-to-end optimization loops. Recent work has shown that conditional diffusion models are strong channel generators, but their inference cost scales with the number of reverse denoising steps. This paper studies conditional drifting models as a low-latency alternative and evaluates them against conditional DDPM, conditional DDIM, and conditional GAN baselines. All models are trained in residual space to approximate the conditional law of $e = y - x$, where x denotes the transmitted symbol and y the channel output. The benchmark covers AWGN, Rayleigh fading, SSPA nonlinearity, and a simplified optical-fiber proxy channel. Our drifting implementation is a practical conditional adaptation of the recently proposed drifting-model idea [1], using an attraction–repulsion drift field on generated residual samples. In the current benchmark, drifting is competitive or best in lower-budget settings and remains strongest on Rayleigh and the optical-fiber proxy at $T = 100$, while DDPM(100) achieves the best sliced Wasserstein distance on AWGN and SSPA. A separate timing study over 20,480 samples shows that one-shot drifting is $260.8\times$ faster than DDPM(100) on CPU. These results position drifting as a promising fast alternative to iterative diffusion-based channel generators.

Index Terms—learned channel models, conditional generative modeling, diffusion models, drifting models, sliced Wasserstein distance, inference latency

I. INTRODUCTION

Accurate channel models are a basic requirement for modern communication-system design. They are especially important in neural and end-to-end learned systems, where the channel appears repeatedly inside training loops, decoder adaptation procedures, and differentiable autoencoder pipelines. When a clean analytical channel model is unavailable, too restrictive, or poorly matched to measured data, generative channel surrogates become attractive because they can learn the full conditional distribution directly from samples.

This problem is not only about fidelity. In many learning pipelines, the channel model is evaluated millions of times. As a result, sample quality, conditioning fidelity, computational cost, and differentiability all matter simultaneously. The question is therefore not merely whether a generative model can match the channel distribution, but whether it can do so at

a cost compatible with large-scale training and deployment workflows.

Recent work on diffusion-based channel modeling has shown that conditional diffusion models can generate high-quality channel samples across standard benchmarks such as AWGN, Rayleigh, and nonlinear amplifier channels [2]. Follow-up work has further studied robustness and speed-quality tradeoffs for diffusion-based channel generation [3], extended diffusion-based channel synthesis to high-dimensional user-specific channels [4], and proposed diffusion-driven digital-twin style generation of statistical channel state information from sensing and location information [5]. More broadly, diffusion methods have also begun to influence adjacent wireless tasks such as high-dimensional channel estimation [6], low-complexity MIMO channel estimation [7], [8], joint channel estimation and data detection [9], and fast time-varying MIMO-OFDM estimation [10], while recent surveys indicate that diffusion models are rapidly becoming a general-purpose tool across wireless networking and communications problems [11], [12]. That line of work is compelling because diffusion models are flexible, robust, and often more stable than earlier adversarial approaches. However, this improvement comes with a practical systems cost: inference requires an iterative reverse denoising procedure, and the latency scales with the number of diffusion steps. This tradeoff is manageable when the channel generator is called occasionally, but it becomes more problematic when it is embedded in inner loops of iterative training or repeated Monte Carlo evaluation.

This motivates the study of one-shot conditional generators for learned channel simulation. Among recently proposed generative paradigms, drifting models [1] are particularly interesting because they aim to learn the transport from a simple latent source to the target data distribution directly in the network parameters, rather than unfold that transport as a multi-step sampling trajectory at inference time. Intuitively, diffusion spends computation at test time by progressively transforming noise into data through a reverse-time chain. Drifting instead spends computation during training: generated samples are repeatedly nudged toward the data manifold by a drift field, and the neural generator is updated so that

future samples move closer to the desired distribution in a single forward pass. Once trained, sampling is therefore naturally one-shot. This perspective is also consistent with recent geometric analyses of generative dynamics that relate diffusion-like behavior to autonomous vector-field structure and marginalized noise geometry [13].

That distinction makes drifting conceptually appealing for channel modeling. In a communication setting, the generator must be conditional on the transmitted symbol x , and it must capture not just a marginal sample cloud, but the conditional law $p(y | x)$. The recently proposed drifting framework is still new, and in the context of communication channels there is no widely adopted baseline yet. This makes it important to document not only benchmark results, but also the concrete conditional adaptation required to make drifting operational in a channel-learning codebase.

The goal of this draft is therefore twofold. First, we extend the existing diffusion-based channel-learning repository of [2] with a practical conditional drifting baseline, a stronger conditional GAN baseline, and an optical-fiber proxy channel adapted from DM_OptFib. Second, we benchmark these models under a common residual-space evaluation protocol using sliced Wasserstein distance (SWD), together with direct timing measurements. The emphasis of this draft is on generator-level distributional fidelity and computational tradeoffs rather than on downstream symbol-error-rate curves.

The main contributions are summarized as follows.

- We formulate a conditional drifting generator for channel simulation by modeling residuals $e = y - x$ and conditioning the generator directly on the transmitted symbol.
- We provide a practical drifting objective based on attraction and repulsion in residual space, which is straightforward to integrate into an existing PyTorch channel-learning pipeline.
- We compare conditional DDPM, conditional DDIM, conditional GAN, and conditional drifting models on AWGN, Rayleigh, SSPA, and optical-fiber proxy channels.
- We quantify the fidelity–latency tradeoff and show that one-shot drifting provides large inference-time savings while remaining competitive in distributional fidelity.

II. PROBLEM FORMULATION AND DRIFTING BACKGROUND

A. What drifting changes conceptually

Because drifting is still unfamiliar to many readers, it is useful to state its generative viewpoint explicitly. Following [1], let f_θ map a latent variable $z \sim p_z$ to a sample $u = f_\theta(z)$, and let

$$q_\theta = (f_\theta)_\# p_z \quad (1)$$

denote the pushforward distribution induced by the generator. The training objective is no longer expressed as a likelihood surrogate or as a denoising score-matching loss. Instead, the paper introduces a *drifting field* that tells us how samples drawn from the current model distribution should move so

that the model distribution approaches the target distribution p .

At a high level, the desired property is an equilibrium condition: if the model already matches the target distribution, then the drift should vanish. Using $V_{p,q}(u)$ to denote the field evaluated at a sample u under target distribution p and model distribution q , the fixed-point intuition can be summarized as

$$q = p \implies V_{p,q}(u) = 0, \quad (2)$$

which is the sense in which matched distributions form an equilibrium of the training dynamics [1], [14].

This makes the training logic very different from diffusion. In diffusion, one learns a denoiser and then performs iterative refinement at inference time. In drifting, one instead constructs improved targets during training by moving current samples along the field $V_{p,q}$, and then regresses the generator onto those improved targets. As a result, the iterative correction happens during optimization rather than at test time.

The one-step view can be written as follows. Starting from a generated sample $u = f_\theta(z)$, define a drifted target

$$\tilde{u} = \text{stopgrad}(u + \eta V_{p,q_\theta}(u)), \quad (3)$$

and optimize the generator by

$$\mathcal{L}_{\text{drift}} = \mathbb{E}_{z \sim p_z} [\|f_\theta(z) - \tilde{u}\|_2^2]. \quad (4)$$

The stop-gradient is essential here: it makes the drifted sample behave like a fixed target for the current optimization step. Repeated training steps therefore evolve the pushforward q_θ toward the target distribution. This is also the cleanest way to see why drifting is naturally one-shot at inference: once the generator has learned to emit samples near equilibrium, sampling only requires drawing z and evaluating $f_\theta(z)$, with no reverse-time chain.

Fig. 3 gives an actual toy particle simulation in the spirit of the particle-based demonstrations used in [1]. For two fixed channel inputs, we initialize mismatched conditional residual particle clouds and then evolve them under the same attraction–repulsion drift rule used in our channel benchmark. The figure therefore illustrates real drift dynamics rather than a hand-drawn schematic, while still serving only as a toy analogue of the full learned generator. To show that the same idea can train an explicit parameterized generator, Fig. 4 uses a tiny conditional MLP trained with the drifting objective on the same toy target. For contrast, Fig. 5 shows an oracle conditional diffusion-style reverse trajectory on the same target, making the difference in sampling path length visually explicit.

B. Conditional channel simulation

We now specialize the generic drifting view to learned conditional channel simulation. Let $x \in \mathbb{R}^n$ denote the transmitted channel input and $y \in \mathbb{R}^n$ the channel output. The objective is to learn a conditional generator that reproduces the channel law given x . In this setting, conditioning enters through the transmitted symbol, while stochasticity is carried by a latent variable $z \sim \mathcal{N}(0, I)$.

TABLE I
TWO TARGET PARAMETERIZATIONS FOR CONDITIONAL DRIFTING

Mode	Output of g_θ	Target object	Short interpretation
Direct $y x$	$\hat{y}$	Learn $p(y x)$ directly	Learns the channel law directly; cleaner and more principled
Residual	$\hat{e}$, then $\hat{y} = x + \hat{e}$	Learn $p(e x)$ with $e = y - x$	Learns the channel perturbation; more biased but often easier to optimize

Two target parameterizations are useful in the experiments. The first is the direct channel-output formulation, in which the generator approximates $p(y | x)$ and predicts $\hat{y} = g_\theta(x, z)$. The second is the residual formulation, in which

$$e = y - x \quad (5)$$

is modeled instead, so that the generator approximates $p(e | x)$ and the channel output is reconstructed as

$$\hat{y} = x + \hat{e}. \quad (6)$$

The direct formulation corresponds to learning the channel law itself, whereas the residual formulation injects the inductive bias that the channel output can be understood as the transmitted symbol plus a structured perturbation. As discussed later in the experiments, the direct formulation is the more natural comparator to direct-output diffusion and WGAN baselines, while the residual formulation is often easier to optimize in the unified benchmark.

C. Conditional drifting model

Our implementation should be viewed as a practical conditional adaptation inspired by [1], rather than as an exhaustive reproduction of every variant in that paper. To avoid confusion with the communication input symbol x , we use u above for the generic sample variable in the abstract formulation and reserve x below for the transmitted symbol. In the residual benchmark, the target object is the residual $e = y - x$, and the generator is therefore

$$\hat{e} = g_\theta(x, z), \quad z \sim \mathcal{N}(0, I), \quad (7)$$

and is implemented as a two-hidden-layer multilayer perceptron with hidden width 128, latent dimension 16, and SiLU activations:

$$[x, z] \rightarrow \text{Linear} \rightarrow \text{SiLU} \rightarrow \text{Linear} \rightarrow \text{SiLU} \rightarrow \text{Linear}. \quad (8)$$

The training target is formed from an attraction–repulsion drift field in residual space. This should be read as a concrete channel-specific instantiation of the generic field $V_{p,q}$. Given generated samples $G = \{g_i\}$ and positive residual samples $P = \{p_j\}$, we compute

$$V(g_i) = \left(\sum_j w_{ij}^{(P)} p_j - g_i \right) - \lambda \left(\sum_k w_{ik}^{(G)} g_k - g_i \right), \quad (9)$$

where $w^{(P)}$ and $w^{(G)}$ are row-normalized RBF kernel weights. The first term attracts each generated sample toward a local

barycenter of true residual samples, and the second term repels it from nearby generated samples to reduce collapse. This yields a detached target

$$\tilde{g}_i = \text{stopgrad}(g_i + \eta V(g_i)), \quad (10)$$

and the generator is optimized by

$$\mathcal{L}_{\text{drift}} = \frac{1}{N} \sum_i \|g_i - \tilde{g}_i\|_2^2. \quad (11)$$

In other words, the generic update $u \mapsto u + \eta V_{p,q}(u)$ from the drifting formulation is realized here with a kernel-based residual drift field estimated directly from minibatches of true and generated residual samples. In the corrected benchmark, the drift scale is $\eta = 1.0$, the repulsive weight is $\lambda = 1.0$, the minimum kernel bandwidth is 0.2, and the maximum drift norm is 2.0. These settings were chosen because they produced stable residual distributions and materially improved the early AWGN drifting baseline.

The crucial operational consequence is straightforward: once training is complete, sampling requires only a single evaluation of $g_\theta(x, z)$. In residual mode, this produces a channel residual sample $\hat{e}$ followed by reconstruction $\hat{y} = x + \hat{e}$; in direct mode, it produces $\hat{y}$ directly. There is no reverse chain, no ODE solver, and no per-sample iterative schedule at inference.

III. EXPERIMENTAL SETUP

A. Channel models

We benchmark four channel families. Unless stated otherwise, each sample is represented as a real-valued I/Q vector

$$x = \begin{bmatrix} x_I \\ x_Q \end{bmatrix} \in \mathbb{R}^2, \quad y = \begin{bmatrix} y_I \\ y_Q \end{bmatrix} \in \mathbb{R}^2. \quad (12)$$

Thus the unified benchmark uses channel dimension $n = 2$, which corresponds to one complex channel input-output pair written in real coordinates. Noise variables are sampled independently across real components unless a channel-specific transformation introduces coupling.

AWGN:

$$y = x + n, \quad n \sim \mathcal{N}(0, \sigma^2 I). \quad (13)$$

This is the standard additive white Gaussian noise reference channel and serves as the simplest near-identity baseline.

Rayleigh: the transmitted symbol is multiplied elementwise by a random fading amplitude and then corrupted by additive Gaussian noise. In the implemented benchmark,

$$y = h \odot x + n, \quad (14)$$

where $\odot$ denotes elementwise multiplication,

$$h_k = \frac{1}{\sqrt{2}} \sqrt{u_k^2 + v_k^2}, \quad u_k, v_k \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1), \quad (15)$$

and

$$n \sim \mathcal{N}(0, \sigma^2 I). \quad (16)$$

Each real component is therefore scaled by an independent Rayleigh-distributed magnitude before noise is added. This

Training step for fixed channel input x





Fig. 1. Training-time drifting update for conditional channel residual modeling. For a fixed transmitted symbol x , the generator produces residual samples $\hat{e}_i$, which are moved toward the true residual cloud through a drift field that combines attraction to data and repulsion among generated samples. The network is then regressed toward the drifted targets.

Residual-mode one-shot inference



After training: draw z , evaluate $g_\theta(x, z)$ once, and reconstruct the channel output as $\hat{y} = x + \hat{e}$.

Constant-depth,
one-shot sampling.

Fig. 2. Inference-time schematic for residual-mode conditional drifting. Once the generator has been trained, sampling requires only a single forward pass of $g_\theta(x, z)$, followed by residual-to-output reconstruction. In the direct-output parameterization, the same one-shot generator would produce $\hat{y}$ directly instead of $\hat{e}$.

channel is multiplicative rather than additive and is correspondingly less aligned with a pure residual perturbation view.

SSPA: the input is passed through a smooth solid-state power amplifier nonlinearity before additive noise is applied. Let

$$r = \sqrt{x_I^2 + x_Q^2} \quad (17)$$

denote the input amplitude. The nonlinear amplitude law used in the benchmark is

$$a(r) = \frac{G}{\left(1 + \left(\frac{Gr}{A_0}\right)^{2p}\right)^{1/(2p)}}, \quad (18)$$

with default parameters $p = 3$, $A_0 = 1.5$, and $G = 5.0$. The noiseless amplifier output is

$$\tilde{x} = a(r)x, \quad (19)$$

and the observed channel output is

$$y = \tilde{x} + n, \quad n \sim \mathcal{N}\left(0, \frac{\sigma^2}{2}I\right). \quad (20)$$

This channel introduces amplitude-dependent nonlinear distortion while preserving the input direction in the I/Q plane.

OptFib: we include a simplified memoryless optical-fiber proxy channel adapted from DM_OptFib. Let the state after




Fig. 3. Toy conditional particle simulation in residual space. Two different channel inputs induce two different target residual distributions. Starting from mismatched generated particle clouds, the middle panel applies one attraction–repulsion drift update of the form $\tilde{e} = \text{stopgrad}(\hat{e} + \eta V(\hat{e}))$, and the last panel shows the particle clouds after repeated updates. This figure is an actual toy simulation of the drift field, not a hand-drawn schematic.




Fig. 4. Toy conditional MLP trained with the drifting objective. The left panel shows target conditional residual samples. The middle panel shows samples from a randomly initialized tiny conditional MLP $g_\theta(x, z)$. The right panel shows the same MLP after training with the drift objective, where the generated samples move close to the correct condition-dependent residual clouds.

propagation step k be $x^{(k)} = [x_I^{(k)}, x_Q^{(k)}]^\top$. At each step, the instantaneous nonlinear phase is

$$\phi^{(k)} = \gamma \left(\left(x_I^{(k)} \right)^2 + \left(x_Q^{(k)} \right)^2 \right) \frac{L}{K_{\text{step}}}. \quad (21)$$

The deterministic Kerr-style rotation is

$$\begin{aligned} \tilde{x}_I^{(k+1)} &= x_I^{(k)} \cos \phi^{(k)} - x_Q^{(k)} \sin \phi^{(k)}, \\ \tilde{x}_Q^{(k+1)} &= x_I^{(k)} \sin \phi^{(k)} + x_Q^{(k)} \cos \phi^{(k)}, \end{aligned} \quad (22)$$

followed by additive Gaussian perturbations at every step:

$$\begin{aligned} x_I^{(k+1)} &= \tilde{x}_I^{(k+1)} + \sigma_n \xi_I^{(k)}, \\ x_Q^{(k+1)} &= \tilde{x}_Q^{(k+1)} + \sigma_n \xi_Q^{(k)}, \end{aligned} \quad \xi_I^{(k)}, \xi_Q^{(k)} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1). \quad (23)$$

When the optical-noise parameterization of DM_OptFib is used,

$$\sigma_n = \sqrt{\frac{10(P_n - 30)/10}{2K_{\text{step}}}}, \quad (24)$$

with P_n measured in dBm. In the current benchmark, the default parameters are $L = 5000$, $\gamma = 1.27$, $K_{\text{step}} = 50$, and $P_n = -21.3$ dBm. This is the most nonlinear channel in the suite and serves as a stress test for all generators.

B. Training protocol

All models are trained conditionally on x in residual space. The local benchmark configuration used for the stored comparison figures is:

- residual dimension $n = 2$,
- noise standard deviation 0.3,
- training set size 30,000,
- diffusion epochs 20,
- GAN epochs 60,
- batch size 256,
- evaluation set size 5,000,
- hidden width 128 for diffusion, drifting, and GAN base-lines,




Fig. 5. Oracle toy conditional diffusion trajectory on the same Gaussian conditional target. The panels show reverse-time snapshots from a 100-step diffusion process at $t = 100$, $t = 50$, $t = 20$, and $t = 0$. Starting from condition-independent noise, the reverse denoising trajectory progressively separates the conditions and contracts toward their target residual distributions. This figure is included to contrast the explicit multi-step reverse path of diffusion with the one-shot sampling path of drifting.

- latent dimension 16 for drifting and GAN generators.

For diffusion, we evaluate two horizons: $T = 20$ and $T = 100$. At each horizon, we report full DDPM sampling, full DDIM sampling, and a reduced-step DDIM trajectory as an ablation. For readability, the fast DDIM ablation is omitted from the main bar plot, but it is retained in the tables because it is informative about quality degradation under aggressive trajectory shortening.

In addition to the local benchmark, the repository includes a dedicated publication runner, `publication_benchmark_t100.py`, for multi-seed GPU evaluation. That script uses a larger configuration with dataset size 120,000, diffusion epochs 60, GAN epochs 120, batch size 512, evaluation size 20,000, bootstrap confidence intervals, and pairwise win tables.

C. Evaluation metric and sanity checks

The primary metric is sliced Wasserstein distance on residual vectors:

$$\text{SWD}(\{\hat{e}_i\}, \{e_i\}). \quad (25)$$

Given true residual samples $P = \{e_i\}_{i=1}^N$ and generated samples $Q = \{\hat{e}_i\}_{i=1}^N$, SWD is estimated by drawing random unit directions θ_m , projecting both clouds onto those directions, computing the one-dimensional Wasserstein distance for each projection, and averaging over projections:

$$\text{SWD}(P, Q) = \frac{1}{M} \sum_{m=1}^M W_1(\{\langle \theta_m, e_i \rangle\}_{i=1}^N, \{\langle \theta_m, \hat{e}_i \rangle\}_{i=1}^N). \quad (26)$$

Thus SWD compares the generated and true sample clouds through many one-dimensional views and captures geometric mismatch more faithfully than a purely bin-based marginal comparison. In the current code, SWD is estimated using 256 random one-dimensional projections. Residual-space SWD was selected over histogram L_1 distance because it better captures geometry in $\mathbb{R}^2$ and behaves more consistently across different channel families.

Several sanity checks were added to support trust in the benchmark pipeline. These include SWD identity and symmetry tests, byte-identical reruns under fixed seeds, finite-value validation of stored JSON result files, and small cross-seed ranking checks on AWGN and OptFib. These checks do not replace the full multi-seed publication run, but they reduce the risk that the observed comparisons are artifacts of an inconsistent implementation.

IV. RESULTS

A. Lower-budget comparison at $T = 20$

Table II shows the lower-budget benchmark. At this shorter horizon, drifting is already highly competitive and attains the best SWD on AWGN, Rayleigh, and OptFib. The improved GAN baseline is also reasonably strong on SSPA. The main takeaway is that when the diffusion budget is constrained, one-shot methods become particularly attractive.

B. Main comparison at $T = 100$

Table III reports the higher-fidelity comparison with $T = 100$. This is the fairest setting for diffusion because it allows the DDPM baseline to use the longer reverse trajectory that it benefits from. The results show a more nuanced picture than a simple “drifting versus diffusion” narrative. DDPM(100) is strongest on AWGN and SSPA, which suggests that iterative refinement remains advantageous on near-Gaussian and moderately nonlinear channels. However, drifting remains the best model on Rayleigh and OptFib in the current benchmark, indicating that the one-shot generator is not merely a speed-oriented compromise.

The DDIM results are also informative. Full DDIM trajectories are consistently weaker than DDPM in the present implementation, and aggressive fast-DDIM sampling degrades sharply on Rayleigh and OptFib. This is precisely why the reduced-step DDIM bars were removed from the main figure: they dominate the plot scale without changing the qualitative conclusion.

TABLE II
RESIDUAL SWD COMPARISON AT $T = 20$ (LOWER IS BETTER)

Channel	DDPM(20)	DDIM(20)	DDIM(10)	GAN	Drifting
AWGN	0.0317	0.0979	0.1407	0.0678	0.0150
Rayleigh	0.0597	0.2668	0.4296	0.1553	0.0219
SSPA	0.0366	0.1020	0.1679	0.0267	0.0191
OptFib	0.1727	0.2176	0.4101	0.2963	0.0632

TABLE III
RESIDUAL SWD COMPARISON AT $T = 100$ (LOWER IS BETTER)

Channel	DDPM(100)	DDIM(100)	DDIM(20)	GAN	Drifting
AWGN	0.0055	0.0623	0.2848	0.0713	0.0154
Rayleigh	0.0484	0.1490	1.0641	0.1459	0.0192
SSPA	0.0180	0.1018	0.5833	0.0324	0.0197
OptFib	0.1662	0.2754	2.2907	0.2957	0.0808




Fig. 6. Residual SWD comparison across channels for the $T = 100$ benchmark. The fast reduced-step DDIM ablation is omitted for readability.

C. Example residual distributions

Fig. 7 provides a channel-level illustration on AWGN. The residual histograms show that DDPM and DDIM capture the overall shape of the residual law, while the drifting model produces a comparably concentrated one-shot residual distribution. The corresponding residual-SWD summary makes the ranking easier to read than a scatter plot alone. More importantly, this figure demonstrates why a geometry-aware metric is useful: the visual overlap between methods can be deceptively similar even when the transport-based discrepancy differs.

D. Inference-time comparison

The central practical motivation for drifting is inference cost. To isolate latency from fidelity, we timed DDPM(100) and drifting separately on CPU over 20,480 generated samples per repeat, using batch size 512, 40 batches, and 7 repeats.

DDPM(100) required 4.2612 s per repeat, corresponding to 4,806 samples/s or 0.2081 ms per sample. The one-shot drifting generator required 0.0163 s per repeat, corresponding to 1.253×10^6 samples/s or 7.98×10^{-4} ms per sample. This yields an observed speedup of $260.8\times$ in favor of drifting.

The timing result is consistent with the model structures in Fig. 2. Diffusion allocates substantial test-time computation to iterative denoising, whereas drifting allocates that work to parameter learning and uses constant-depth sampling once training is complete. In workloads where the channel model is invoked repeatedly, this distinction can dominate overall runtime.

E. Interpretation and limitations

Several caveats are important. First, the tables above are current single-seed benchmark results. They are useful for method development and paper drafting, but the final pub-




Fig. 7. AWGN example with training curves, residual distributions, sample clouds, and residual SWD comparison.

lication version should be based on the multi-seed $T = 100$ runner included in the repository. Second, the present study focuses on generator-level distributional fidelity and timing rather than on full end-to-end communication metrics such as SER or E_b/N_0 curves. Relative to [2], the present study should therefore be interpreted as a complementary benchmark on channel-distribution learning rather than as a replacement for downstream system evaluation.

Third, the drifting implementation here is deliberately pragmatic. It is inspired by [1] and captures the core one-shot training idea through an attraction–repulsion drift field, but it is not a claim that every modeling choice from the drifting paper has been reproduced exactly. That said, the current results are already strong enough to justify a more systematic follow-up study.

V. CONCLUSION

This paper presented a detailed experimental draft for conditional drifting models in learned channel simulation and compared them against conditional diffusion and GAN baselines on AWGN, Rayleigh, SSPA, and optical-fiber proxy channels. The results show a clear tradeoff. At lower diffusion budgets, drifting is highly competitive and often strongest. At $T = 100$, DDPM remains the best fidelity baseline on AWGN and SSPA, while drifting is strongest on Rayleigh and OptFib and preserves a one-shot inference path. Most importantly, drifting provides a $260.8\times$ CPU speedup over DDPM(100) in the measured configuration.

These observations suggest that drifting is not only an interesting new generative-model idea, but also a practically relevant direction for communication-system channel learning. The next steps are straightforward: run the full multi-seed publication benchmark, report confidence intervals, and extend the comparison to downstream communication metrics and more realistic optical or memory channels.

REFERENCES

- [1] M. Deng, H. Li, T. Li, Y. Du, and K. He, “Generative modeling via drifting,” *arXiv preprint arXiv:2602.04770*, 2026. [Online]. Available: <https://arxiv.org/abs/2602.04770>
- [2] M. Kim, R. Fritschek, and R. F. Schaefer, “Diffusion models for accurate channel distribution generation,” *arXiv preprint arXiv:2309.10505*, 2023. [Online]. Available: <https://arxiv.org/abs/2309.10505>
- [3] —, “Robust generation of channel distributions with diffusion models,” in *IEEE International Conference on Communications (ICC)*, 2024. [Online]. Available: <https://portal.fis.tum.de/en/publications/robust-generation-of-channel-distributions-with-diffusion-models>
- [4] Y. Lee, J. Choi, B. Kim, and S. Bahk, “Generating high dimensional user-specific wireless channels using diffusion models,” *arXiv preprint arXiv:2409.03924*, 2024. [Online]. Available: <https://arxiv.org/abs/2409.03924>
- [5] C. Gong, F. Yang, S. Yan, H. Zhang, M. Di Renzo, S. J. Wang, and M.-S. Alouini, “Digital twin of channel: Diffusion model for sensing-assisted statistical channel state information generation,” *IEEE Transactions on Wireless Communications*, 2025. [Online]. Available: <https://pure.bit.edu.cn/en/publications/digital-twin-of-channel-diffusion-model-for-sensing-assisted-stat>
- [6] Y. Zhou, Y. Yan, J. Lei, P. Tang, H. Song, and Z. Han, “Generative diffusion models for high dimensional channel estimation,” *IEEE Transactions on Wireless Communications*, 2025. [Online]. Available: <https://dblp.org/rec/journals/corr/abs-2408-10501.html>
- [7] B. Fesl, W. Utschick, H. Wymeersch, and S. Aïssa, “Diffusion-based generative prior for low-complexity mimo channel estimation,” *IEEE Wireless Communications Letters*, 2024. [Online]. Available: https://github.com/benediktfesl/Diffusion_channel_est
- [8] S. Chen, W. Shin, and A. Nallanathan, “Generative diffusion model-based variational inference for mimo channel estimation,” *IEEE Transactions on Communications*, 2025. [Online]. Available: <https://khu.elsevierpure.com/en/publications/generative-diffusion-model-based-variational-inference-for-mimo-c>
- [9] N. Zilberstein, A. Swami, and S. Segarra, “Joint channel estimation and data detection in massive mimo systems based on diffusion models,” in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2024. [Online]. Available: <https://www.catalyzex.com/paper/joint-channel-estimation-and-data-detection>
- [10] Y. Liu, Y. Zhang, Z. Wang, and J. Zhang, “Conditional denoising diffusion-based channel estimation for fast time-varying mimo-ofdm systems,” *Digital Signal Processing*, 2025. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S1051200425003057>
- [11] H. Fan, J. Chen, G. Li, W. Xiao, Y. Wang, and H. V. Poor, “Generative diffusion models for wireless networks: Fundamental, architecture, and state-of-the-art,” *arXiv preprint arXiv:2507.16733*, 2025. [Online]. Available: <https://arxiv.org/abs/2507.16733>
- [12] D.-T. Luong, M. Assi, A. Ghayeb, and M. Debbah, “Diffusion models for future networks and communications: A comprehensive survey,” *arXiv preprint arXiv:2508.01586*, 2025. [Online]. Available: <https://arxiv.org/abs/2508.01586>
- [13] S. Sahraee-Ardakan, M. Delbracio, and P. Milanfar, “The geometry of noise: Why diffusion models don’t need noise conditioning,” *arXiv preprint arXiv:2602.18428*, 2026. [Online]. Available: <https://arxiv.org/abs/2602.18428>
- [14] M. Deng, H. Li, T. Li, Y. Du, and K. He, “Generative modeling via drifting: Project page,” <https://lambertae.github.io/projects/drifting/>, 2026, accessed: 2026-03-16.